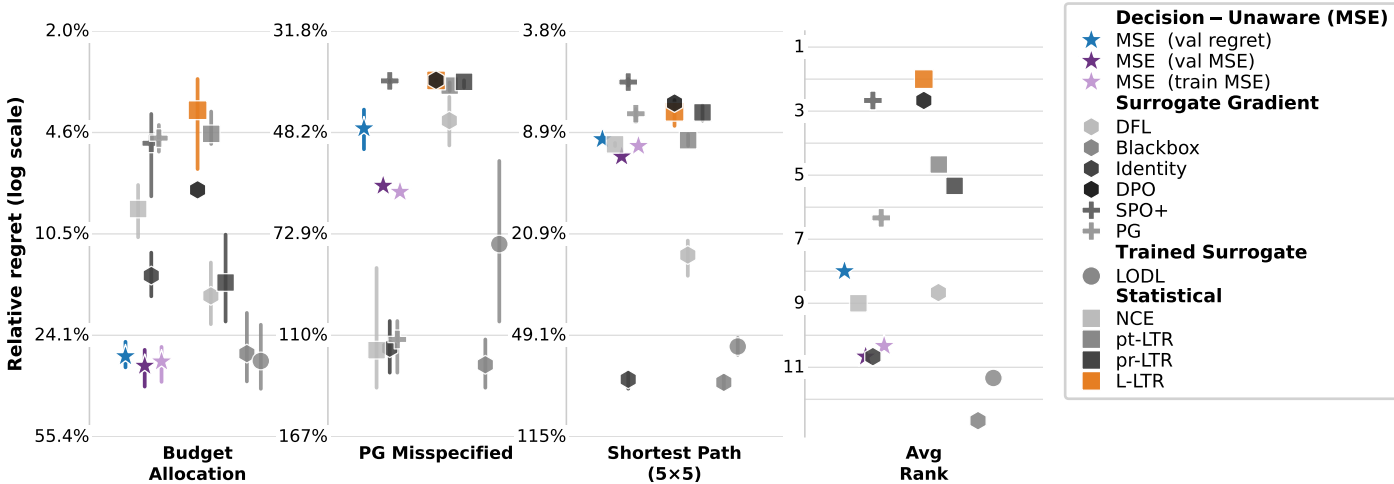
Benchmark re-run — 10-seed mean ± 95% bootstrap CI — Relative regret % (best ↑)
Relative regret (log scale)
2.0%
4.6%
10.5%
24.1%
55.4%
31.8%
48.2%
72.9%
110%
167%
3.8%
8.9%
20.9%
49.1%
115%
1
3
5
7
9
11
Budget Allocation
PG Misspecified
Shortest Path (5×5)
Avg Rank
Decision – Unaware (MSE)
MSE (val regret)
MSE (val MSE)
MSE (train MSE)
Surrogate Gradient
DFL
Blackbox
Identity
DPO
SPO+
PG
Trained Surrogate
LODL
Statistical
NCE
pt-LTR
pr-LTR
L-LTR